

Assignments

Operation	Comment	Call	Team
Example			
$B := LB$	Lower triangular matrix L	<code>Trmm_llnn (L, B)</code>	Robert vdG
Symmetric matrix-matrix multiplication			
$C := AB + C$	A symmetric, stored in lower-triangular part	<code>Symm_ll (A, B, C)</code>	Ben Nguyen, Andrew Dong-Tran, Jeffrey Leung
$C := AB + C$	A symmetric, stored in upper-triangular part	<code>Symm_lu (A, B, C)</code>	Wesley Chung, Hiep Vu, David Shi
$C := BA + C$	A symmetric, stored in lower-triangular part	<code>Symm_rl (A, B, C)</code>	Jeff Taube and Justin Salazar and Darya Mylius
$C := BA + C$	A symmetric, stored in upper-triangular part	<code>Symm_ru (A, B, C)</code>	Liangkun Zhao, Yajie Niu
Symmetric rank-k update			
$C := AA^T + C$	C symmetric, stored in lower-triangular part	<code>Syrk_lu (A, C)</code>	Jorge Munoz, Max Svetlik
$C := AA^T + C$	C symmetric, stored in upper-triangular part	<code>Syrk_ut (A, C)</code>	
$C := A^T A + C$	C symmetric, stored in lower-triangular part	<code>Syrk_lu (A, C)</code>	
$C := A^T A + C$	C symmetric, stored in upper-triangular part	<code>Syrk_ut (A, C)</code>	

Operation	Comment	Team	
Symmetric rank-2k update			
$C := AB^T + BA^T + C$	C symmetric, stored in lower-triangular part	Syr2k_ln (A, B, C)	Ryan Young and Benny Renard
$C := AB^T + BA^T + C$	C symmetric, stored in upper-triangular part	Syr2k_un (A, B, C)	Scott Munro and Matthew Chin
$C := A^T B + B^T A + C$	C symmetric, stored in lower-triangular part	Syr2k_lt (A, B, C)	Raeeca Narimani and Allison Wallpole
$C := A^T B + B^T A + C$	C symmetric, stored in upper-triangular part	Syr2k_ut (A, B, C)	Pushkar and Sarah
Triangular matrix-matrix multiplication			
$B := L^T B$	L stored in lower triangle	Trmm_lltn (A, B)	Sean Wang and Marco Guajardo, Tim, Vincent
$B := UB$	U stored in upper triangle	Trmm_lunn (A, B)	
$B := U^T B$	U stored in upper triangle	Trmm_lutn (A, B)	Chris Getz, Bradley Holloway, Rohit Pattanaik
$B := BL$	L stored in lower triangle	Trmm_rlnn (A, B)	Ben Yang, Andras Balogh, Eric Lee
$B := BL^T$	L stored in lower triangle	Trmm_rltn (A, B)	Amanda Nguyen, Tim Kwan
$B := BU$	U stored in upper triangle	Trmm_runn (A, B)	
$B := BU^T$	U stored in upper triangle	Trmm_rutn (A, B)	
Triangular-triangular matrix multiplication			
$B := LU$	L lower triangular, U upper triangular	Trtrmm_lunn (L, U, B)	
$B := UL$	L lower triangular, U upper triangular	Trtrmm_ulnn (U, L, B)	
$B := U^T L$	L lower triangular, U upper triangular	Trtrmm_ultn (U, L, B)	
$B := UL^T$	L lower triangular, U upper triangular	Trtrmm_ult (U, L, B)	

Part I

Triangular Matrix-matrix Multiplication

Cases

Triangular matrix-matrix multiplication (TRMM) is matrix-matrix multiplication where one of the matrices is square and (lower or upper) triangular.

The full set of triangular matrix-matrix multiplication cases is denoted by TRMM_□□□□, where the letters in the four boxes denote whether the triangular matrix is on the Left or Right, is Lower or Upper triangular, is Not transposed or Transposed, and has a Nonunit or Unit diagonal:

	LL□□	LU□□	RL□□	RU□□
□□N□	$B = LB$	$B = UB$	$B = BL$	$B = BU$
□□T□	$B = L^T B$	$B = U^T B$	$B = BL^T$	$B = BU^T$

$B := LB$ — Team: Robert van de Geijn

Consider the operation

$$B := LB$$

where L is a $m \times m$ lower triangular matrix and B is a $m \times n$ matrix. This is a special case of triangular matrix-matrix multiplication, with the Lower triangular matrix on the LEFT, and the triangular matrix is Not transposed. We will refer to this operation as TRMM_LLNN where the LLNN stands for left lower no-transpose nonunit diagonal. The nonunit diagonal means we will use the entries of the matrix that are stored on the diagonal.

2.1 Precondition and postcondition

In the precondition

$$B = \hat{B}$$

\hat{B} denotes the original contents of B . This allows us to express the state upon completion, the postcondition, as

$$B = L\hat{B}.$$

2.2 Partitioned Matrix Expressions and loop invariants

There are two PME's for this operation.

2.2.1 PME 1

To derive the second PME, partition

$$L \rightarrow \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right), \quad \text{and} \quad B \rightarrow \left(\begin{array}{c} B_T \\ B_B \end{array} \right).$$

Substituting these into the postcondition yields

$$\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} L_{TL} & 0 \\ L_{BL} & L_{BR} \end{pmatrix} \begin{pmatrix} \hat{B}_T \\ \hat{B}_B \end{pmatrix}$$

or, equivalently,

$$\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} L_{TL}\hat{B}_T \\ L_{BL}\hat{B}_T + L_{BR}\hat{B}_B \end{pmatrix}$$

so that, upon completion

$$\frac{B_T = L_{TL}\hat{B}_T}{B_B = L_{BL}\hat{B}_T + L_{BR}\hat{B}_B}$$

From this, we can choose two loop invariants:

Invariant 1: $\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} \hat{B}_T \\ L_{BR}\hat{B}_B \end{pmatrix}$. (The top part has been left alone and the bottom part has been partially computed).

Invariant 2: $\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} \hat{B}_T \\ L_{BL}\hat{B}_T + L_{BR}\hat{B}_B \end{pmatrix}$. (The top part has been left alone and the bottom part has been completely computed).

2.2.2 PME 2

To derive the second PME, partition

$$B \rightarrow \left(B_L \mid B_R \right)$$

and does not partition L . Substituting these into the postcondition yields

$$\left(B_L \mid B_R \right) = L \left(\hat{B}_L \mid \hat{B}_R \right)$$

or, equivalently,

$$\left(B_L \mid B_R \right) = \left(L\hat{B}_L \mid L\hat{B}_R \right)$$

so that, upon completion

$$B_L = L\hat{B}_L \mid B_R = L\hat{B}_R$$

From this, we can choose two more loop invariants:

Invariant 3: $\left(B_L \mid B_R \right) = \left(L\hat{B}_L \mid \hat{B}_R \right)$. (The left part has been completely finished and the right part has been left untouched).

Invariant 4: $\left(B_L \mid B_R \right) = \left(\hat{B}_L \mid L\hat{B}_R \right)$. (The left part has been completely finished and the right part has been left untouched).

2.2.3 Notes

How do I decide to partition the matrices in the postcondition?

- Pick a matrix (operand), any matrix.
- If that matrix has
 - a triangular structure (in storage), then you want to either partition it into four quadrants, or not at all. Symmetric matrices and triangular matrices have a triangular structure (in storage).
 - no particular structure, then you partition it vertically (left-right), horizontally (top-bottom), or not at all.
- Next, partition the other matrices similarly, but conformally (meaning the resulting multiplications with the parts are legal).

Take our problem here: $B := LB$. Start by partitioning L into quadrants:

$$L = \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \hat{B}.$$

Now, the way partitioned matrix multiplication works, this doesn't make sense:

$$B = \underbrace{\left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \hat{B}}_{\left(\begin{array}{c} L_{TL} \times \text{something} + 0 \times \text{something} \\ \hline L_{BL} \times \text{something} + L_{BR} \times \text{something} \end{array} \right)}.$$

So, we need to also partition B into a top part and a bottom part:

$$\left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right) = \underbrace{\left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \left(\begin{array}{c} \hat{B}_T \\ \hline \hat{B}_B \end{array} \right)}_{\left(\begin{array}{c} L_{TL}B_T + B_B \\ \hline L_{BL}B_T + L_{BR}B_B \end{array} \right)}.$$

Alternatively, what if you don't partition L ? You have to partition *something* so let's try partitioning B :

$$\left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right) = L \left(\begin{array}{c} \hat{B}_T \\ \hline \hat{B}_B \end{array} \right)$$

But that doesn't work... Instead

$$\left(\begin{array}{c|c} B_L & B_R \end{array} \right) = L \left(\begin{array}{c|c} \hat{B}_L & \hat{B}_R \end{array} \right) = \left(\begin{array}{c|c} L\hat{B}_L & L\hat{B}_R \end{array} \right)$$

works just fine.

2.3 Determining the Loop Guard

Let us pick one of the loop invariants, Invariant 2, for our further discussion:

$$\left(\frac{B_T = \hat{B}_T}{B_B = L_{BL}\hat{B}_T + L_{BR}\hat{B}_B} \right)$$

We fill this in in the worksheet as Step 2.

To determine the loop guard, you look at this condition and the postcondition and you ask yourself the question “Under what condition, G , does

$$\left(\frac{B_T}{B_B} \right) = \left(\frac{\hat{B}_T}{L_{BL}\hat{B}_T + L_{BR}\hat{B}_B} \right) \wedge \neg G$$

imply the postcondition

$$B = L\hat{B}?$$

The answer is “When $\neg G$ implies that B_B is all of B and (hence) L_{BR} is all of L .” So, we take G to be $m(B_B) < m(B)$ (or $m(L_{BR}) < m(L)$). This is Step 3 in the worksheet.

2.4 Determining the initialization

Next, we turn to Step 4, the initialization. To determine it, we consider the precondition, the partitioning into submatrices of L and B , and the loop invariant we choose.

1	$\{B = \hat{B}\}$
4	$L \rightarrow \left(\begin{array}{c c} L_{TL} & L_{TR} \\ \hline L_{BL} & L_{BR} \end{array} \right), B \rightarrow \left(\begin{array}{c} B_T \\ B_B \end{array} \right)$ where L_{BR} is $? \times ?$, B_B has ? rows
2	$\left\{ \left(\frac{B_T}{B_B} \right) = \left(\frac{\hat{B}_T}{L_{BL}\hat{B}_T + L_{BR}\hat{B}_B} \right) \right\}$

The question now is “What choice of size ? will make this code segment correct?” The answer is to choose ? as 0.

2.5 Progressing Through the Matrices

Next, we must determine how to “march” through the matrices. Notice that initially L_{BR} and B_B are 0×0 and $0 \times n$, respectively. Eventually, L_{BR} must be all of L and B_B all of B . This suggests that we march through matrix L from bottom-right to top-left and through B from bottom to top.

2.5.1 Unblocked algorithm

For the unblocked algorithm, this translates to

5a	$\left(\begin{array}{c c} L_{TL} & L_{TR} \\ \hline L_{BL} & L_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c c c} L_{00} & l_{01} & L_{02} \\ \hline l_{10}^T & \lambda_{11} & l_{12}^T \\ \hline L_{20} & l_{21} & L_{22} \end{array} \right), \left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right) \rightarrow \left(\begin{array}{c} B_0 \\ \hline b_1^T \\ \hline B_2 \end{array} \right)$ <p>where λ_{11} is 1×1, b_1 has 1 row</p>
\vdots	\vdots
5b	$\left(\begin{array}{c c} L_{TL} & L_{TR} \\ \hline L_{BL} & L_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c c c} L_{00} & l_{01} & L_{02} \\ \hline l_{10}^T & \lambda_{11} & l_{12}^T \\ \hline L_{20} & l_{21} & L_{22} \end{array} \right), \left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right) \leftarrow \left(\begin{array}{c} B_0 \\ \hline b_1^T \\ \hline B_2 \end{array} \right)$

Notice that the special shape of L is not of particular interest at this point. The “partition-repartition-move on” merely indicates how we are moving through the array L .

2.5.2 Blocked algorithm

For the Blocked algorithm, this instead translates to

5a	$\left(\begin{array}{c c} L_{TL} & L_{TR} \\ \hline L_{BL} & L_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c c c} L_{00} & L_{01} & L_{02} \\ \hline L_{10} & L_{11} & L_{12} \\ \hline L_{20} & L_{21} & L_{22} \end{array} \right), \left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right) \rightarrow \left(\begin{array}{c} B_0 \\ \hline B_1 \\ \hline B_2 \end{array} \right)$ <p>where L_{11} is $b \times b$, B_1 has b rows</p>
\vdots	\vdots
5b	$\left(\begin{array}{c c} L_{TL} & L_{TR} \\ \hline L_{BL} & L_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c c c} L_{00} & L_{01} & L_{02} \\ \hline L_{10} & L_{11} & L_{12} \\ \hline L_{20} & L_{21} & L_{22} \end{array} \right), \left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right) \leftarrow \left(\begin{array}{c} B_0 \\ \hline B_1 \\ \hline B_2 \end{array} \right)$

Again, the special shape of L is not of particular interest at this point. The “partition-repartition-move on” merely indicates how we are moving through the array L .

2.6 State Before the Update

Next, we determine what the state of the exposed submatrices of B is, in terms of the exposed submatrices of L .

2.6.1 Unblocked algorithm

Start with the loop invariant:

$$\left(\frac{B_T}{B_B} \right) = \left(\frac{\hat{B}_T}{L_{BL}\hat{B}_T + L_{BR}\hat{B}_B} \right)$$

By examining Step 5a we find that

$$\begin{array}{c|c|c} L_{TL} \rightarrow \left(\begin{array}{c|c} L_{00} & l_{01} \\ \hline l_{10}^T & L_{11} \end{array} \right) & L_{TR} \rightarrow \left(\begin{array}{c} L_{02} \\ \hline l_{12}^T \end{array} \right) & \text{and} \quad B_T \rightarrow \left(\begin{array}{c} B_0 \\ \hline b_1^T \end{array} \right) \\ \hline L_{BL} \rightarrow \left(\begin{array}{c|c} L_{20} & l_{21} \end{array} \right) & L_{BR} \rightarrow L_{22} & B_B \rightarrow B_2 \end{array}$$

Substituting this into the loop invariant yields

$$\left(\frac{\left(\begin{array}{c} B_0 \\ \hline b_1^T \end{array} \right)}{B_2} \right) = \left(\frac{\left(\begin{array}{c} \hat{B}_0 \\ \hline \hat{b}_1^T \end{array} \right)}{\left(\begin{array}{c|c} L_{20} & l_{21} \end{array} \right) \left(\begin{array}{c} \hat{B}_0 \\ \hline \hat{b}_1^T \end{array} \right) + L_{22}\hat{B}_2} \right)$$

or, applying the rules of linear algebra and partitioned matrix-matrix multiplication,

$$\left(\frac{\begin{array}{c} B_0 \\ \hline b_1^T \\ \hline B_2 \end{array}} \right) = \left(\frac{\begin{array}{c} \hat{B}_0 \\ \hline \hat{b}_1^T \end{array}}{L_{20}\hat{B}_0 + l_{21}\hat{b}_1^T + L_{22}\hat{B}_2} \right)$$

2.6.2 Blocked algorithm

Again, start with the loop invariant:

$$\left(\frac{B_T}{B_B} \right) = \left(\frac{\hat{B}_T}{L_{BL}\hat{B}_T + L_{BR}\hat{B}_B} \right)$$

By examining Step 5a we find that

$$\begin{array}{c|c|c} L_{TL} \rightarrow \left(\begin{array}{c|c} L_{00} & L_{01} \\ \hline L_{10} & L_{11} \end{array} \right) & L_{TR} \rightarrow \left(\begin{array}{c} L_{02} \\ \hline L_{12} \end{array} \right) & \text{and} \quad B_T \rightarrow \left(\begin{array}{c} B_0 \\ \hline B_1 \end{array} \right) \\ \hline L_{BL} \rightarrow \left(\begin{array}{c|c} L_{20} & L_{21} \end{array} \right) & L_{BR} \rightarrow L_{22} & B_B \rightarrow B_2 \end{array}$$

Substituting this into the loop invariant yields

$$\left(\frac{\left(\begin{array}{c} B_0 \\ \hline B_1 \end{array} \right)}{B_2} \right) = \left(\frac{\left(\begin{array}{c} \hat{B}_0 \\ \hline \hat{B}_1 \end{array} \right)}{\left(\begin{array}{c|c} L_{20} & L_{21} \end{array} \right) \left(\begin{array}{c} \hat{B}_0 \\ \hline \hat{B}_1 \end{array} \right) + L_{22}\hat{B}_2} \right)$$

or, applying the rules of linear algebra and partitioned matrix-matrix multiplication,

$$\begin{pmatrix} B_0 \\ B_1 \\ B_2 \end{pmatrix} = \begin{pmatrix} \hat{B}_0 \\ \hat{B}_1 \\ L_{20}\hat{B}_0 + L_{21}\hat{B}_1 + L_{22}\hat{B}_2 \end{pmatrix}$$

2.7 State After the Update

Next, we determine what the state of the exposed submatrices of B is, in terms of the exposed submatrices of L *after* the update (Step 7).

2.7.1 Unblocked algorithm

Again, start with the loop invariant:

$$\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} \hat{B}_T \\ L_{BL}\hat{B}_T + L_{BR}\hat{B}_B \end{pmatrix}$$

By examining Step 5b we find that

$$\begin{array}{c|c} L_{TL} \rightarrow L_{00} & L_{TR} \rightarrow \left(\begin{array}{c|c} L_{02} & l_{12}^T \end{array} \right) \\ \hline L_{BL} \rightarrow \left(\begin{array}{c} l_{10}^T \\ L_{20} \end{array} \right) & L_{BR} \rightarrow \left(\begin{array}{c|c} \lambda_{11} & l_{12}^T \\ l_{21} & L_{22} \end{array} \right) \end{array} \quad \text{and} \quad \begin{array}{c} B_T \rightarrow B_0 \\ \hline B_B \rightarrow \left(\begin{array}{c} b_1^T \\ B_2 \end{array} \right) \end{array}$$

Substituting this into the loop invariant yields

$$\begin{pmatrix} B_0 \\ \left(\begin{array}{c} b_1^T \\ B_2 \end{array} \right) \end{pmatrix} = \begin{pmatrix} \hat{B}_0 \\ \left(\begin{array}{c} l_{10}^T \\ L_{20} \end{array} \right) \hat{B}_0 + \left(\begin{array}{c|c} \lambda_{11} & 0 \\ l_{21} & L_{22} \end{array} \right) \left(\begin{array}{c} \hat{b}_1^T \\ \hat{B}_2 \end{array} \right) \end{pmatrix}.$$

Notice that we *now* take into account that L is lower triangular and hence $l_{12}^T = 0$! Applying the rules of linear algebra and partitioned matrix-matrix multiplication yields

$$\begin{pmatrix} B_0 \\ b_1^T \\ B_2 \end{pmatrix} = \begin{pmatrix} \hat{B}_0 \\ \frac{l_{12}^T \hat{B}_0 + \lambda_{11} \hat{b}_1^T}{L_{20}\hat{B}_0 + l_{21}\hat{b}_1^T + L_{22}\hat{B}_2} \end{pmatrix}$$

2.7.2 Blocked algorithm

Again, start with the loop invariant:

$$\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} \hat{B}_T \\ L_{BL}\hat{B}_T + L_{BR}\hat{B}_B \end{pmatrix}$$

By examining Step 5b we find that

$$\begin{array}{c|c} L_{TL} \rightarrow L_{00} & L_{TR} \rightarrow \left(\begin{array}{c|c} L_{02} & L_{12} \end{array} \right) \\ \hline L_{BL} \rightarrow \left(\begin{array}{c} L_{10} \\ L_{20} \end{array} \right) & L_{BR} \rightarrow \left(\begin{array}{c|c} L_{11} & L_{12} \\ \hline l_{21} & L_{22} \end{array} \right) \end{array} \quad \text{and} \quad \begin{array}{c} B_T \rightarrow B_0 \\ \hline B_B \rightarrow \left(\begin{array}{c} B_1 \\ B_2 \end{array} \right) \end{array}$$

Substituting this into the loop invariant yields

$$\left(\begin{array}{c} B_0 \\ \hline \left(\begin{array}{c} B_1 \\ B_2 \end{array} \right) \end{array} \right) = \left(\begin{array}{c} \hat{B}_0 \\ \hline \left(\begin{array}{c} L_{10} \\ L_{20} \end{array} \right) \hat{B}_0 + \left(\begin{array}{c|c} L_{11} & 0 \\ \hline L_{21} & L_{22} \end{array} \right) \left(\begin{array}{c} \hat{B}_1 \\ \hline \hat{B}_2 \end{array} \right) \end{array} \right).$$

Notice that we *now* take into account that L is lower triangular and hence $L_{12} = 0$! Applying the rules of linear algebra and partitioned matrix-matrix multiplication yields

$$\left(\begin{array}{c} B_0 \\ \hline B_1 \\ B_2 \end{array} \right) = \left(\begin{array}{c} \hat{B}_0 \\ \hline \frac{L_{12}\hat{B}_0 + L_{11}\hat{B}_1}{L_{20}\hat{B}_0 + L_{21}\hat{B}_1 + L_{22}\hat{B}_2} \end{array} \right)$$

2.8 Determining the update

2.8.1 Unblocked algorithm

Now we are ready to determine the update, by comparing Step 6 and Step 7:

6	$\left\{ \left(\begin{array}{c} B_0 \\ \hline b_1^T \\ B_2 \end{array} \right) = \left(\begin{array}{c} \hat{B}_0 \\ \hline \frac{\hat{b}_1^T}{L_{20}\hat{B}_0 + l_{21}\hat{b}_1^T + L_{22}\hat{B}_2} \end{array} \right) \right\}$
8	
7	$\left\{ \left(\begin{array}{c} B_0 \\ \hline b_1^T \\ B_2 \end{array} \right) = \left(\begin{array}{c} \hat{B}_0 \\ \hline \frac{l_{12}^T \hat{B}_0 + \lambda_{11} \hat{b}_1^T}{L_{20}\hat{B}_0 + l_{21}\hat{b}_1^T + L_{22}\hat{B}_2} \end{array} \right) \right\}$

The update now becomes

$$\begin{aligned} b_1^T &:= \lambda_{11} b_1^T \\ b_1^T &:= l_{12}^T B_0 + b_1^T \end{aligned}$$

We break this up into two steps, because of the kind of routines we have at our disposal. The first is a call to `FLA_Scal` while the second is a call to `FLA_Gemv`. Notice that the order here matters!

2.8.2 Blocked algorithm

To determine the update, by comparing Step 6 and Step 7:

6	$\left\{ \left(\begin{array}{c} B_0 \\ B_1 \\ B_2 \end{array} \right) = \left(\begin{array}{c} \hat{B}_0 \\ \hat{B}_1 \\ \frac{L_{20}\hat{B}_0 + L_{21}\hat{B}_1 + L_{22}\hat{B}_2}{L_{20}\hat{B}_0 + L_{21}\hat{B}_1 + L_{22}\hat{B}_2} \end{array} \right) \right\}$
8	
7	$\left\{ \left(\begin{array}{c} B_0 \\ B_1 \\ B_2 \end{array} \right) = \left(\begin{array}{c} \hat{B}_0 \\ \frac{L_{12}\hat{B}_0 + L_{11}\hat{B}_1}{L_{20}\hat{B}_0 + L_{21}\hat{B}_1 + L_{22}\hat{B}_2} \\ \frac{L_{20}\hat{B}_0 + L_{21}\hat{B}_1 + L_{22}\hat{B}_2}{L_{20}\hat{B}_0 + L_{21}\hat{B}_1 + L_{22}\hat{B}_2} \end{array} \right) \right\}$

The update now becomes

$$B_1 := L_{11}B_1$$

$$B_1 := L_{12}B_0 + B_1$$

We break this up into two steps, because of the kind of routines we have at our disposal. The first is a call to `Trmm_llnn_unb_var1` (notice that L_{11} is lower triangular, so you have to be careful) while the second is a call to `FLA_Gemm`. Notice that the order here matters!

Part II

Symmetric Matrix-matrix Multiplication

Cases

Symmetric matrix-matrix multiplication (SYMM) is matrix-matrix multiplication where one of the matrices is square and symmetric, stored in the lower or upper triangular part of the matrix.

The full set of symmetric matrix-matrix multiplication cases is denoted by SYMM_□□, where the letters in the two boxes denote whether the symmetric matrix is on the Left or Right, and is stored in the Lower or Upper triangular part of that matrix:

	L□	R□	
□L	$C := AB + C$	$C := BA + C$	A symmetric, stored in lower triangle
□U	$C := AB + C$	$C := BA + C$	A symmetric, stored in upper triangle

Part III

Symmetric Rank-k Update

Cases

Symmetric rank-k update (SYRK) is matrix-matrix multiplication of a matrix with its transpose, where the matrix C being updated is symmetric, stored in the lower or upper triangular part of the matrix.

The full set of symmetric rank-k cases is denoted by SYRK_□□, where the letters in the two boxes denote whether the matrices begin multiplied are Not transposed or Transposed, and whether the matrix being updated is stored in the Lower or Upper triangular part:

	N□	T□	
□L	$C := AA^T + C$	$C := A^T A + C$	C symmetric, stored in lower triangle
□U	$C := AA^T + C$	$C := A^T A + C$	C symmetric, stored in upper triangle

Part IV

Symmetric Rank-2k Update

Cases

Symmetric rank-2k update (SYR2K) is matrix-matrix multiplication of two matrices where both the result and the transpose of that result are added to a symmetric matrix C , stored in the lower or upper triangular part of the matrix.

The full set of symmetric rank-2k cases is denoted by SYR2K_□□, where the letters in the two boxes denote whether the matrices begin multiplied are Not transposed or Transposed, and whether the matrix being updated is stored in the Lower or Upper triangular part:

	N□	T□	
□L	$C := AB^T + BA^T + C$	$C := A^T B + B^T A + C$	C symmetric, stored in lower triangle
□U	$C := AB^T BA^T + C$	$C := A^T B + B^T A + C$	C symmetric, stored in upper triangle

Part V

Triangular Solve with Multiple Right-hand Sides

Cases

Triangular solve with multiple right-hand sides (TRSM) is a matrix-matrix multiplication in disguise: Not only is one of the matrices square and (lower or upper) triangular, but in addition, one multiplies with the inverse of the matrix.

The full set of TRSM cases is denoted by `TRSM_□□□□`, where the letters in the four boxes denote whether the triangular matrix is on the Left or Right, is Lower or Upper triangular, is Not transposed or Transposed, and has a Nonunit or Unit diagonal:

	LL□□	LU□□	RL□□	RU□□
□□N□	$B := L^{-1}B$	$B := U^{-1}B$	$B := BL^{-1}$	$B := BU^{-1}$
□□T□	$B := L^{-T}B$	$B := U^{-T}B$	$B := BL^{-T}$	$B := BU^{-T}$

We will only consider the case where we compute with the diagonal.

Where does the name *triangular solve with multiple right-hand sides* come from? If one wants to compute $B := A^{-1}B$ one can also think of this as solving $AX = B$, overwriting the matrix A with the solution matrix X . Now, partition X and B by columns. Then

$$\underbrace{A \begin{pmatrix} x_0 & x_1 & \cdots & x_{n-1} \end{pmatrix}}_{\begin{pmatrix} Ax_0 & Ax_1 & \cdots & Ax_{n-1} \end{pmatrix}} = \begin{pmatrix} b_0 & b_1 & \cdots & b_{n-1} \end{pmatrix}$$

and hence, for each column j , $Ax_j = b_j$. This means that one can instead solve with matrix A , and because B has many columns, this becomes a *solve with multiple right-hand sides*. If A is triangular, then it is a *triangular solve with multiple right-hand sides*. Importantly, the inverse of the matrix is never computed.